

# Work Assignment

## Phase 3

João Pedro Dias Faria

Dept. de informática - UC de Computação Paralela  
Universidade do Minho  
Braga, Portugal  
pg53939

Rui Pedro Guise da Silva

Dept. de informática - UC de Computação Paralela  
Universidade do Minho  
Braga, Portugal  
pg54213

**Abstract**—his document serves as an auxiliary support for the Parallel Computing course project, which corresponds to the academic year 2022/2023.

**Index Terms**—C, SeARCH, OpenMP, CUDA, Parallelization.

### I. INTRODUCTION

In the concluding phase of our project, we have undertaken an ambitious initiative to enhance the execution time of our previously introduced algorithm. Our primary objective is to optimize the program further by implementing more efficient parallelization techniques, such as increasing the number of available threads without compromising on overhead or distributing algorithmic tasks among processes.

This endeavor aligns with our broader strategy, as our exploration extends to the domain of Graphics Processing Units (GPUs), renowned for their proficiency in parallel computations. GPUs create an environment where large, uninterrupted data blocks are efficiently handled, allowing numerous threads to operate simultaneously with minimal overhead. To harness this potential, we are incorporating CUDA (Compute Unified Device Architecture) into our development process.

Furthermore, this report serves as a detailed chronicle of our journey through these complex enhancements. It comprehensively details the technical challenges encountered, the strategic implementations pursued, and the resulting improvements in program performance.

### II. WA1

The initial stage of this practical assignment involved implementing optimization techniques to minimize execution time. These techniques encompassed alterations to algorithms, enhancements in ILP (Instruction-Level Parallelism) and vectorization. Through complexity analysis and profiling, we identified two functions that consumed nearly 100% of the execution time: 'Potential()' with 54.92% of the total computational effort and 'computeAccelerations()', considering that it is executed approximately 201 times by the 'VelocityVerlet()' function, being responsible for 41.57% of the program's execution time.




Fig. 1. Results from gprof2dot for the non-optimised code

In this initial phase, we simplified mathematical expressions to decrease the number of operations and modified the algorithms used. The most significant change in this first phase was the merging of the two functions into a single one, named 'computeAccelerationsOPT'.

### III. WA2

In the second phase, after optimizing the sequential code, we explored parallelization techniques using OpenMP. To identify the function that takes up the most execution time, we used perf record on the sequential code from the first phase. As anticipated, the most time-consuming function is 'computeAccelerationsOPT,' and it was this function that applied parallelization. In summary, these were the changes made during this phase. We used **#pragma omp parallel for** on the outer loop so that each thread is responsible for a set of particles. All other variables are declared within this loop to remain private to each thread. Additionally, we used reduction for the variables 'a' and 'P', where a copy of the 'P' and 'a' variables is made for each thread, and the values of these copies are aggregated into the original variables at the end of the loop. This addition was made because these variables are used outside the loop. The last change made in this phase was to add **schedule(dynamic, 40)** since the threads responsible for the last particles had less workload compared to the first ones.

#### IV. TESTS AND METRICS TO BE EVALUATED

To assess the efficiency of our algorithm, we will employ two key performance metrics: execution time, which measures the total duration of algorithm execution, and Level 1 cache misses, which quantify the number of times data is retrieved from slower memory due to its absence in the CPU's primary cache. These metrics are strongly correlated with the overall performance improvement of an application, as reducing execution time and minimizing cache misses directly expedite processing and enhance overall system responsiveness.

The tests will be performed two computing platforms: the SEARCH Cluster, which has been provided by *Universidade do Minho* and has been extensively utilized in the earlier phases of this project, along with a personal laptop equipped with the following specifications:

Lenovo Legion	
OS	Linux Mint 21.2
CPU	Intel(R) Core(TM) i7-10750H CPU @ 2.60GHz
RAM	16GB
GPU	NVIDIA GeForce GTX 1650 Ti

For each set of conditions, our team will execute 3 tests and calculate the average of the metrics, to further reduce variances inside the same machine.

#### V. KEY BLOCKS

# Overhead	Samples	Command	Shared Object	Symbol
#	#			
90.96%	1627	MDseq.exe	MDseq.exe	[.] computeAccelerationsOPT
9.01%	154	MDseq.exe	[unknown]	[k] 0xffffffffa158cbc0
0.01%	4	MDseq.exe	[unknown]	[k] 0xffffffffa0e6d5f8
0.01%	1	MDseq.exe	[unknown]	[k] 0xffffffffa0ee4b4f
0.00%	1	MDseq.exe	[unknown]	[k] 0xffffffffa1026a73
0.00%	1	MDseq.exe	libc-2.17.so	[.] __random
0.00%	1	MDseq.exe	[unknown]	[k] 0xffffffffa158c48a
0.00%	1	MDseq.exe	MDseq.exe	[.] initializeVelocities
0.00%	1	MDseq.exe	ld-2.17.so	[.] check_match.9525
0.00%	1	MDseq.exe	[unknown]	[k] 0xffffffffa0ffb653
0.00%	1	MDseq.exe	[unknown]	[k] 0xffffffffa159608e
0.00%	1	MDseq.exe	[unknown]	[k] 0xffffffffa0fa9edc
0.00%	1	MDseq.exe	ld-2.17.so	[.] __libc_memalign
0.00%	1	MDseq.exe	[unknown]	[k] 0xffffffffa1005afc
0.00%	1	MDseq.exe	[unknown]	[k] 0xffffffffa1213df1
0.00%	1	MDseq.exe	[unknown]	[k] 0xffffffffa0eea540
0.00%	3	perf	[unknown]	[k] 0xffffffffa0e6d5f8

Fig. 2. Output obtained using *perf record*

By analyzing the computationally intensive functions in our previous sequential version, it becomes evident that 'computeAccelerationsOPT' is the predominant function of this algorithm, responsible for a significant portion of the execution time.

Through the parallelization of this function utilizing OpenMP (Open Multi-Processing), we achieved a notable reduction in the algorithm's execution time. However, it persists as the most computationally intensive component within our application, indicating that the potential for substantial execution time enhancement lies in the pursuit of more advanced parallelization techniques.

#### VI. FIRST VERSION DEVELOPED

##### A. Strategy

For our first version, we began by implementing CUDA only in the heaviest-weight functions, namely 'computeAccelerations' and 'computeAccelerationsOPT'. For this implementation, we needed to allocate arrays for the particle positions and accelerations, 'r' and 'a' respectively, as well as a variable 'P' related to the value of 'Pot'. Thus, the new variables 'da', 'dr', and 'dP' emerged. This accounts for all the necessary memory for this implementation and is carried out in the 'initializeKernel' function. In the 'initializeKernel' function, we took the opportunity to copy the particle positions to the device. We decided not to copy the array of accelerations since its initialization would be performed by a kernel.

After the kernel initialization, we focused on implementing 'computeAccelerations' with CUDA. In the OMP version, the 'a' matrix was initialized at the beginning of the function 'computeAccelerations', and then its values were updated. Using CUDA, we couldn't perform this entire function with just one kernel call because we couldn't guarantee synchronization across all blocks. Thus, we couldn't ensure that all values were initialized before being updated. The solution found was to split the function into two kernel calls, one for initialization and another for updating, resulting in the functions 'initializeacelarationKernel' and 'newacelarationKernel'. This solution was found after some research and is commonly referred to as 'CPU synchronization'. This synchronization is performed by the host, which waits for all threads to finish before calling the next kernel.

We also use CUDA to assist whenever the 'VelocityVerlet' function wants to call the 'computeAccelerationsOPT' function.

This was replaced by two kernel calls, continuing with the same thought presented earlier, one is the same 'initializeacelarationKernel', and another to update their accelerations and calculate the new value of 'dP', 'newpotacelarationKernel'. Before these two calls, we need to copy the updated positions of the particles to the device, and after the calls, copy the new updates and the value of 'Pot' back to the host.

In the 'newacelarationKernel' and 'newpotacelarationKernel' function, each thread represents a particle and calculates the difference in acceleration that its iteration with the particles in front of it in the array will cause. Since two threads can modify the acceleration of the same particle, we use a 'myatomicAdd' function to write to the 'da' array. The purpose of using this function is to avoid data races, which were prevented in the previous phase version by the reduction. In the 'newpotacelarationKernel', we also use this atomic function to add the value to 'dP'.

##### B. Results

VI		
Test	TExec (s)	L1 Cache Misses
N=5000, SEARCH	7.666	-
N=5000, Laptop	3.208	7280797

A more in-depth analysis of the obtained results will be conducted in one of the final chapters.

### C. Profiling

```
==392== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	99.32%	5.79893s	201	28.811ms	28.778ms	28.852ms	newpotacelarationKernel(int, d
	0.52%	30.441ms	1	30.441ms	30.441ms	30.441ms	newacelarationKernel(int, d
	0.07%	4.2183ms	202	20.882us	20.576us	29.953us	[CUDA memcp HtoD]
	0.07%	3.9046ms	403	9.6880us	1.2790us	18.785us	[CUDA memcp DtoH]
	0.02%	903.24us	202	4.4710us	4.0960us	5.1840us	initializeacelarationKernel
API calls:	96.47%	5.84729s	605	9.6649ms	39.008us	30.559ms	cudaMemcpy
	3.43%	207.95ms	3	69.318ms	7.7110us	207.94ms	cudaMalloc
	0.07%	4.0627ms	404	10.056us	6.1720us	82.136us	cudaLaunchKernel
	0.01%	581.57us	1	581.57us	581.57us	581.57us	cuDeviceTotalMem
	0.01%	338.98us	101	3.3560us	211ns	133.52us	cuDeviceGetAttribute
	0.00%	277.72us	809	343ns	194ns	15.555us	cudaGetLastError
	0.00%	233.90us	3	77.967us	4.7030us	171.88us	cudaFree
	0.00%	223.51us	1	223.51us	223.51us	223.51us	cuDeviceGetName
	0.00%	9.0250us	1	9.0250us	9.0250us	9.0250us	cuDeviceGetPCIBusId
	0.00%	1.5080us	3	502ns	222ns	829ns	cuDeviceGetCount
	0.00%	1.2260us	1	1.2260us	1.2260us	1.2260us	cuDeviceGetUuid
	0.00%	1.0370us	2	518ns	209ns	828ns	cuDeviceGet

Fig. 3. Profiling - version 1

From this analysis, we can identify relevant aspects of our implementation.

- The `cudaMalloc` API call has a small overhead, taking around 207.95 to process.
- The `'initializeacelarationKernel'` was call 202 times, averaging 4.47 us per call, taking a total of 903 us to execute.
- The `'newpotacelarationKernel'` takes 5.79 seconds to complete the 201 calls, averaging 28.811 ms per call. This elevated value is due to the use of `'myatomicAdd'`, which hinders the parallel approach.
- Similary, the `'newacelarationKernel'` takes 30ms in its single call.

## VII. SECOND VERSION DEVELOPED

### A. Strategy

The use of our version of `'atomicAdd'` for double adds significant overhead to our implementation. Therefore, in this new version, our main goal was to eliminate dependency on this function to calculate the new accelerations and continue without data races in our program.

With this goal in mind, we realized that, to avoid using `'atomicAdd'`, each thread could only make changes to the acceleration of the particle ('da') it corresponds to. Thus, a change in the algorithm was necessary, so that instead of each particle calculating its new acceleration and that of the particle it interacts with, those in front in the array, each particle now only calculates its acceleration resulting from its interaction with all the existing particles. With this new implementation, each thread modifies different positions in the array, since they only change the positions of the corresponding particles, eliminating the need for the use of atomic add. The function `'newpotacelarationKernel'` was renamed to `'newpotacelarationKernel1Part'` with this new implementation.

To handle the `'atomicAdd'` of the `'Pot'` variable, instead of allocating space for a single double, we started allocating

space for each existing particle. In the `'newpotacelarationKernel1Part'` function, each thread adds the calculated potential from the iterations of its particle relative to those in front of it, in a position of the new array. This array is copied to the host, which adds all the values to determine the total potential after the new iteration.

With this new implementation, all dependencies on `'atomicAdd'` were removed, achieving the goal of not relying on this function for our implementation, resulting in an improvement in the execution time.

### B. Results

V2		
Test	TExec (s)	L1 Cache Misses
N=5000, SEARCH	5.127	-
N=5000, Laptop	3.199	7754848

We can see that our search execution time had an improvement from 7.666 seconds to 5.127 seconds

### C. Profiling

```
==25226== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	98.71%	1.12382s	201	5.5871ms	5.5752ms	5.6041ms	newpotacelarationKernel1Part
	0.43%	4.8818ms	403	12.113us	6.0480us	18.592us	[CUDA memcp DtoH]
	0.40%	4.5149ms	1	4.5149ms	4.5149ms	4.5149ms	newacelarationKernel(int, d
	0.38%	4.3663ms	202	21.615us	21.280us	22.336us	[CUDA memcp HtoD]
	0.08%	871.20us	202	4.3120us	4.1920us	4.8320us	initializeacelarationKernel
API calls:	82.41%	1.16280s	605	1.9220ms	47.972us	5.6590ms	cudaMemcpy
	17.03%	240.34ms	3	80.112ms	8.5290us	240.32ms	cudaMalloc
	0.36%	5.0280ms	404	12.445us	6.8800us	75.971us	cudaLaunchKernel
	0.00%	1.1139ms	2	556.95us	518.13us	595.77us	cuDeviceTotalMem
	0.07%	972.36us	202	4.8130us	238ns	207.68us	cuDeviceGetAttribute
	0.02%	264.84us	2	132.42us	59.882us	204.95us	cuDeviceGetName
	0.02%	240.27us	809	297ns	192ns	7.9050us	cudaGetLastError
	0.01%	184.64us	3	61.546us	4.1110us	168.45us	cudaFree
	0.00%	19.042us	2	9.5210us	3.9450us	15.097us	cuDeviceGetPCIBusId
	0.00%	3.6730us	4	918ns	394ns	1.4040us	cuDeviceGet
	0.00%	2.5940us	3	864ns	429ns	1.4130us	cuDeviceGetCount
	0.00%	1.6800us	2	840ns	590ns	1.0900us	cuDeviceGetUuid

Fig. 4. Profiling - version 2

- The `cudaMalloc` API take around 240.34ms to process. This increase is due to the fact that we allocate space for more values.
- The `'initializeacelarationKernel'` average 4.31 us per call, taking a total of 871 us to execute.
- `'newpotacelarationKernel1Part'` takes now 1.12 seconds to complete the 201 calls, averaging 5.5871 ms per call. So, we can see that removing the atomics resulted in an improvement of approximately 5.17 times, for this function.
- This improvement was also observed in the `'newacelarationKernel'` function, which now takes 4.5 ms in its call.
- The CUDA memcp from device to host and from host to device took 4.88 ms and 4.36 ms, respectively.

## VIII. THIRD VERSION DEVELOPED

### A. Strategy

In this latest implementation, our goal is to avoid constant copying between the host and device and vice versa, and to

make a better implementation in the calculation of 'dP', using shared memory.

To avoid constant copying between the host and device, we need all arrays related to particles to be updated within the device, so that it always has their values updated. To achieve this, we need to allocate space for the velocity array and start executing the 'VelocityVerlet', 'MeanSquaredVelocity', and 'Kinetic' functions on the device.

The 'VelocityVerlet' function was divided into several functions for reasons similar to those mentioned in the implementation of the 'newacelarationKernel' and 'newpotacelarationKernel'. These new functions include the 'posvelKernel', 'updatevelKernel' and 'elasticKernel', in addition to those already created, 'initializeacelarationKernel' and 'newpotacelarationKernel1Part'.

While implementing these new functions, we faced the same challenge that 'dP' had caused us, as we needed to reduce the values calculated by the threads to a single value. Since one of the objectives of this latest implementation was to take advantage of shared memory in the calculation of 'dP', we used the same solution for each of the problems.

This solution involves each block having a shared array of size equal to the number of threads per block used. Instead of entering its value into the global array, each thread inserts the value into the shared memory array at the position equivalent to its ID in the block. We synchronize the threads of the block using the `__synchronize()` function to ensure that all values are in the array. After this synchronization, we can perform a parallel reduction, with the total value calculated by the block residing in the first position of the array. The first thread of each block writes the obtained result to the global array, which is initialized not with a size of N but with the number of blocks used, in this implementation.

This implementation is used in the function 'elasticKernel', 'MeanSquaredVelocitykernel', 'Kinetickernel', each one with a specific global array being 'dPress', 'dmvs' and 'dKE' respectively. The space for these arrays is allocated during the 'initializeKernel' and is influenced by the number of blocks to be used. Following the same reasoning, 'dP' now allocates space only for the number of blocks instead of the number of particles.

The aggregate value of all these global arrays is calculated by a single thread in the function 'reduce'. The summed value of each one is stored in a specific position of the results array.

Since the values of 'Press', 'mvs', 'KE' and 'P' can be obtained on the host simultaneously, so instead of having a variable for each value, we this 'results' array where each position represents one of the necessary values. To have space on the device to store these results in the 'initializeKernel', we also allocate space for the respective array, as mentioned earlier. This way, we avoid multiple calls to 'memcpy'.

After each reduce call, we copy the 'dresults' array to the host to obtain the values we have just calculated.

## B. Results

Final		
Test	TExec (s)	L1 Cache Misses
N=5000, SEARCH	2.981	-
N=5000, Laptop	3.202	5396658

We can see that our search execution time had an improvement from 5.127 seconds to 2.981 seconds to

## C. Profiling

```
==38383== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	98.41%	1.13956s	281	5.6694ms	5.6553ms	5.6915ms	newpotacelarationKernelPart
	0.44%	5.1088ms	281	25.416us	25.217us	25.921us	reduce(double*, double*, dou
	0.39%	4.5099ms	1	4.5099ms	4.5099ms	4.5099ms	newacelarationKernel(int, do
	0.17%	2.0001ms	281	9.9500us	9.7280us	10.369us	posvelKernel(int, double*, d
	0.15%	1.7397ms	281	8.6550us	8.3200us	9.472us	elasticKernel(int, double*,
	0.12%	1.3561ms	281	6.7460us	6.5920us	6.9760us	MeanSquaredVelocitykernel(in
	0.12%	1.3364ms	281	6.6480us	6.4320us	7.2010us	KineticKernel(int, double*,
	0.11%	1.2392ms	281	6.1650us	6.0160us	6.5920us	updatevelKernel(int, double*
	0.07%	824.98us	282	4.0840us	4.0000us	4.8000us	initializeacelarationKernel(
	0.03%	293.95us	282	1.4550us	1.2800us	18.592us	[CUDA memcpy DtoH]
	0.00%	42.560us	2	21.280us	20.864us	21.696us	[CUDA memcpy HtoD]
API calls:	83.41%	1.15291s	284	5.6515ms	73.102us	5.7346ms	cudaMemcpy
	15.53%	214.70ms	8	26.837ms	3.7090us	214.66ms	cudaMalloc
	1.01%	13.932ms	1618	8.6530us	5.8890us	88.124us	cudaLaunchKernel
	0.62%	282.89us	1	282.89us	282.89us	282.89us	cudaDeviceTotalMem
	0.01%	153.64us	181	1.5210us	149ns	61.346us	cudaDeviceGetAttribute
	0.01%	149.45us	8	18.681us	2.7020us	128.61us	cudaFree
	0.00%	66.937us	286	324ns	220ns	3.6640us	cudaGetLastError
	0.00%	29.650us	1	29.650us	29.650us	29.650us	cudaDeviceGetName
	0.00%	9.1890us	1	9.1890us	9.1890us	9.1890us	cudaDeviceGetPCIBusId
	0.00%	1.6310us	3	543ns	251ns	907ns	cudaDeviceGetCount
	0.00%	997ns	2	498ns	196ns	801ns	cudaDeviceGet
	0.00%	506ns	1	506ns	506ns	506ns	cudaDeviceGetUuid

Fig. 5. Profiling - version 3

- The `cudaMalloc` API take around 214ms to process to allocate all the space necessary. In this way, although more arrays are allocated, the new approach with shared memory helps with the time because we allocate smaller-sized arrays.
- The 'initializeacelarationKernel' average 4.08 us per call, taking a total of 824 us to execute.
- 'newpotacelarationKernel1Part' and 'newacelarationKernel' take approximately the same time as in the last version.
- The 'reduce' average 25.416 us per call, taking a total of 5.1ms.
- The 'posvelKernel' average 9.95 us per call, taking a total of 2ms.
- The 'elasticKernel' average 8.65 us per call, taking a total of 1.73ms.
- The 'MeanSquaredVelocityKernel' average 6.74 us per call, taking a total of 1.3ms.
- The 'KineticKernel' average 6.64 us per call, taking a total of 1.33ms.
- The 'MeanSquaredVelocityKernel' average 6.74 us per call, taking a total of 1.3ms.
- The CUDA memcpy from device to host and from host to device took 293.95 us and 42.56 us, respectively.

## IX. TESTS AND ANALYSIS

### A. General comparison

To better understand the evolution of the execution time throughout this work, we started by presenting the results for a number of particles equal to 5000 ( $N=5000$ ). It was also suggested in this phase to use a personal computer to analyze how our program behaves on another machine. Thus, we once again analyzed how the execution time varies for the various versions and compared the results with those obtained on the search machine. The laptop used has the specifications mentioned earlier. Presenting the execution time for the three versions created in this phase.




Fig. 6. Execution time for the last version ( $N=5000$ )

	Search	Laptop
N=5000 V1	7,666	3,208
N=5000 V2	5,127	3,199
N=5000 V3	2,981	3,202

Fig. 7. Table execution time for the last version ( $N=5000$ )

### B. L1 Cache Misses

We also wanted to explore other metrics in addition to those used in previous phases. Therefore, we decided to analyze the L1 cache misses between the same versions compared earlier. However, we were unable to make a proper comparison with the results obtained on the Search for this type of metric, as we encountered some difficulties in using the metric on the Search cluster. Nevertheless, we decided to retain this metric, confirming by reading the bar chart presented below that our efforts to improve the final program were indeed successful, resulting in a reduction in the number of L1 cache misses from version 2 to version 3.




Fig. 8. L1 cache misses on laptop

	Laptop
N=5000 V1	7280797
N=5000 V2	7754848
N=5000 V3	5396658

Fig. 9. Table L1 cache misses - Laptop

The reduction in L1 cache misses in the last version is related to the fact that more calculations are performed on the GPU. It's normal for the number of cache misses by the CPU to decrease for this version.

### C. Scalability analysis

To better analyze the scalability of our program, we tested how the execution time varies with changes in the number of particles. For each phase, we assessed its scalability based on the value of  $N$  and the required execution time. Once again, we performed this operation both on the Search and on our personal laptop, and obtained the results presented below.

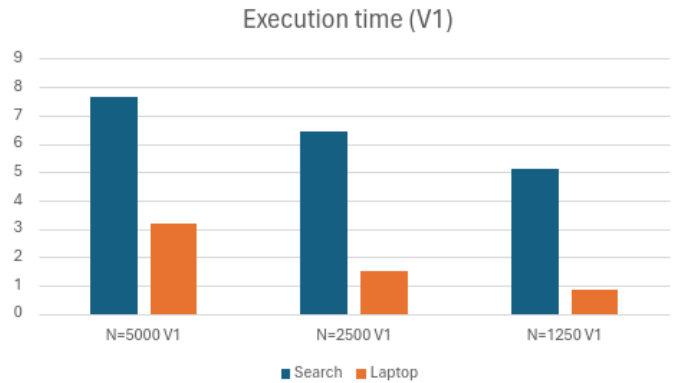


Fig. 10. Execution time - version 1



	Search	Laptop
N=5000 V1	7,666	3,208
N=2500 V1	6,431	1,504
N=1250 V1	5,132	0,87

Fig. 11. Table execution time - version 1




Fig. 12. Execution time - version 2

	Search	Laptop
N=5000 V2	5,127	3,199
N=2500 V2	4,445	1,311
N=1250 V2	3,638	0,735

Fig. 13. Table execution time - version 2




Fig. 14. Execution time - version 3

	Search	Laptop
N=5000 V3	2,981	3,202
N=2500 V3	1,834	1,338
N=1250 V3	1,57	0,833

Fig. 15. Table execution time - version 3

may intensify as the number of particles increases, negatively impacting the ideal scalability.

## X. CONCLUSION

In summary, although the execution time was not better than that obtained in the previous phase, we consider that we made a correct implementation of our problem using CUDA, being aware of the limitations that arose during its execution. This work helped us to better understand how to program using the GPU. This phase also helped to better understand the difficulty of performing parallel reductions and the challenges they pose to programmers in implementing these algorithms.

The reasons for the increase in execution time with the increase in the number of particles may be related to the competition for resources among threads. This competition